

K-Fold Cross Validation & Grid Search & Feature Selection

Importing the libraries

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 import pandas as pd
```

Importing the dataset

```
In [2]: 1 dataset = pd.read_csv('Social_Network_Ads.csv')
        2 X = dataset.iloc[:, :-1]
        3 y = dataset.iloc[:, -1]
```

Splitting the dataset into the Training set and Test set

```
In [3]: 1 from sklearn.model_selection import train_test_split
        2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

Feature Scaling

```
In [4]: 1 from sklearn.preprocessing import StandardScaler
        2 sc = StandardScaler()
        3 X_train = sc.fit_transform(X_train)
        4 X_test = sc.transform(X_test)
```

Training the Kernel SVM model on the Training set

```
In [5]: 1 from sklearn.svm import SVC
        2 classifier = SVC(kernel='rbf', random_state=0)
        3 # classifier = SVC(C=0.25, gamma=0.8, kernel='rbf', random_state=0)
        4 classifier.fit(X_train, y_train)
```

Out[5]: SVC(random_state=0)

Making the Confusion Matrix

```
In [6]: 1 from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
        2
        3 y_pred = classifier.predict(X_test)
        4 print(confusion_matrix(y_test, y_pred))
        5 print(classification_report(y_test, y_pred))
        6 print(accuracy_score(y_test, y_pred))
```

```
[[55  3]
 [ 1 21]]
      precision    recall  f1-score   support

      0       0.98      0.95      0.96         58
      1       0.88      0.95      0.91         22

   accuracy: 0.95
  macro avg: 0.93      0.95      0.94         80
 weighted avg: 0.95      0.95      0.95         80

0.95
```

Applying k-Fold Cross Validation

```
In [7]: 1 from sklearn.model_selection import cross_val_score
        2
        3 accuracies = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=10, verbose=1) # cv: Cross Validation=k=10
        4 print(accuracies)
        5 print(f"Accuracy: {np.round(accuracies.mean()*100, 2)} %") # average accuracy score
        6 print(f"Std: {accuracies.std()*100} %")
```

```
[0.84375 0.875 0.90625 0.84375 0.9375 0.84375 0.90625 0.90625 1.
 0.9375]
Accuracy: 90.0 %
Std: 4.80071609241788 %
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 0.0s finished
```

Applying Grid Search to find the best model and the best parameters

```
In [8]: 1 from sklearn.model_selection import GridSearchCV
2
3 parameters = [{'C': [0.25, 0.5, 0.75, 1], 'kernel': ['linear']},
4               {'C': [0.25, 0.5, 0.75, 1], 'kernel': ['rbf'], 'gamma': [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]}]
5 grid_search = GridSearchCV(estimator = classifier,
6                             param_grid = parameters,
7                             scoring = 'accuracy',
8                             cv = 10, # cv: Cross Validation, k=10
9                             n_jobs = -1, # all cpu
10                            verbose=1 # print
11                            )
12 grid_search.fit(X_train, y_train)
13 print(grid_search.best_estimator_)
14 print(grid_search.best_params_)
15 print(grid_search.best_score_) # precision
16 print(grid_search.best_index_)
```

```
Fitting 10 folds for each of 40 candidates, totalling 400 fits
SVC(C=0.25, gamma=0.8, random_state=0)
{'C': 0.25, 'gamma': 0.8, 'kernel': 'rbf'}
0.90625
11
```

n_jobs设定工作的core数量

等于-1的时候，表示cpu里的所有core进行工作。

▼ grid_search里的最好score

```
In [9]: 1 grid_predictions = grid_search.predict(X_test)
2
3 from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
4
5 print(classification_report(y_test, grid_predictions))
6 print(accuracy_score(y_test, grid_predictions))
7 print(confusion_matrix(y_test, grid_predictions))
```

	precision	recall	f1-score	support
0	0.98	0.95	0.96	58
1	0.88	0.95	0.91	22
accuracy			0.95	80
macro avg	0.93	0.95	0.94	80
weighted avg	0.95	0.95	0.95	80

```
0.95
[[55  3]
 [ 1 21]]
```

▼ 手动调最好score

```
In [10]: 1 classifier = SVC(C=0.25, gamma=0.8, kernel='rbf', random_state=0)
2 classifier.fit(X_train, y_train)
3
4 from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
5
6 predictions = classifier.predict(X_test)
7 print(classification_report(y_test, predictions))
8 print(accuracy_score(y_test, predictions))
9 print(confusion_matrix(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.98	0.95	0.96	58
1	0.88	0.95	0.91	22
accuracy			0.95	80
macro avg	0.93	0.95	0.94	80
weighted avg	0.95	0.95	0.95	80

```
0.95
[[55  3]
 [ 1 21]]
```

▼ Feature Selection 递归特征消除 (Recursive feature elimination) and Re-do Model Training

```
In [11]: 1 from sklearn.feature_selection import RFE # 递归特征消除 (Recursive feature elimination)
2
3 classifier = SVC(C=0.25, gamma=0.8, kernel='rbf', random_state=0)
4 rfe = RFE(estimator=classifier, n_features_to_select=10) # find 10 columns that are most valuable
5 rfe.fit(X_train, y_train)
```

```
Out[11]: RFE(estimator=SVC(C=0.25, gamma=0.8, random_state=0), n_features_to_select=10)
```

```
In [12]: 1 # Two columns should be included
2 rfe.support_
```

```
Out[12]: array([ True,  True])
```

```
In [13]: 1 rfe.ranking_
```

```
Out[13]: array([1, 1])
```

```
In [14]: 1 X = dataset.iloc[:, 1:]
```

411 [47].

```
1 X = dataset.iloc[:, :-1]
2 X = X[X.columns[rfe.support_]]
3 y = dataset.iloc[:, -1]
4
5 from sklearn.model_selection import train_test_split
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
7
8 from sklearn.preprocessing import StandardScaler
9 sc = StandardScaler()
10 X_train = sc.fit_transform(X_train)
11 X_test = sc.transform(X_test)
12
13 from sklearn.svm import SVC
14 classifier = SVC(C=0.25, gamma=0.8, kernel='rbf', random_state=0)
15 classifier.fit(X_train, y_train)
16
17 from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
18
19 predictions = classifier.predict(X_test)
20 print(classification_report(y_test, predictions))
21 print(accuracy_score(y_test, predictions))
22 print(confusion_matrix(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.98	0.95	0.96	58
1	0.88	0.95	0.91	22
accuracy			0.95	80
macro avg	0.93	0.95	0.94	80
weighted avg	0.95	0.95	0.95	80

```
0.95
[[55  3]
 [ 1 21]]
```